

# CS & IT ENGINEERING

## Operating Systems

### System Calls & threads

Lecture No. 1



By- Dr. Khaleel Khan Sir





TOPICS TO BE  
COVERED


**System Call Implementation**

**Fork System Call**



# System Calls

is a means of availing o.s services;



```
main()  
{
```

um { f(); user defined  
scanf(); Pre defined

k.m → fork(); OS defined

↓  
Privileged Instrn / Slw Instr.  
Instr.



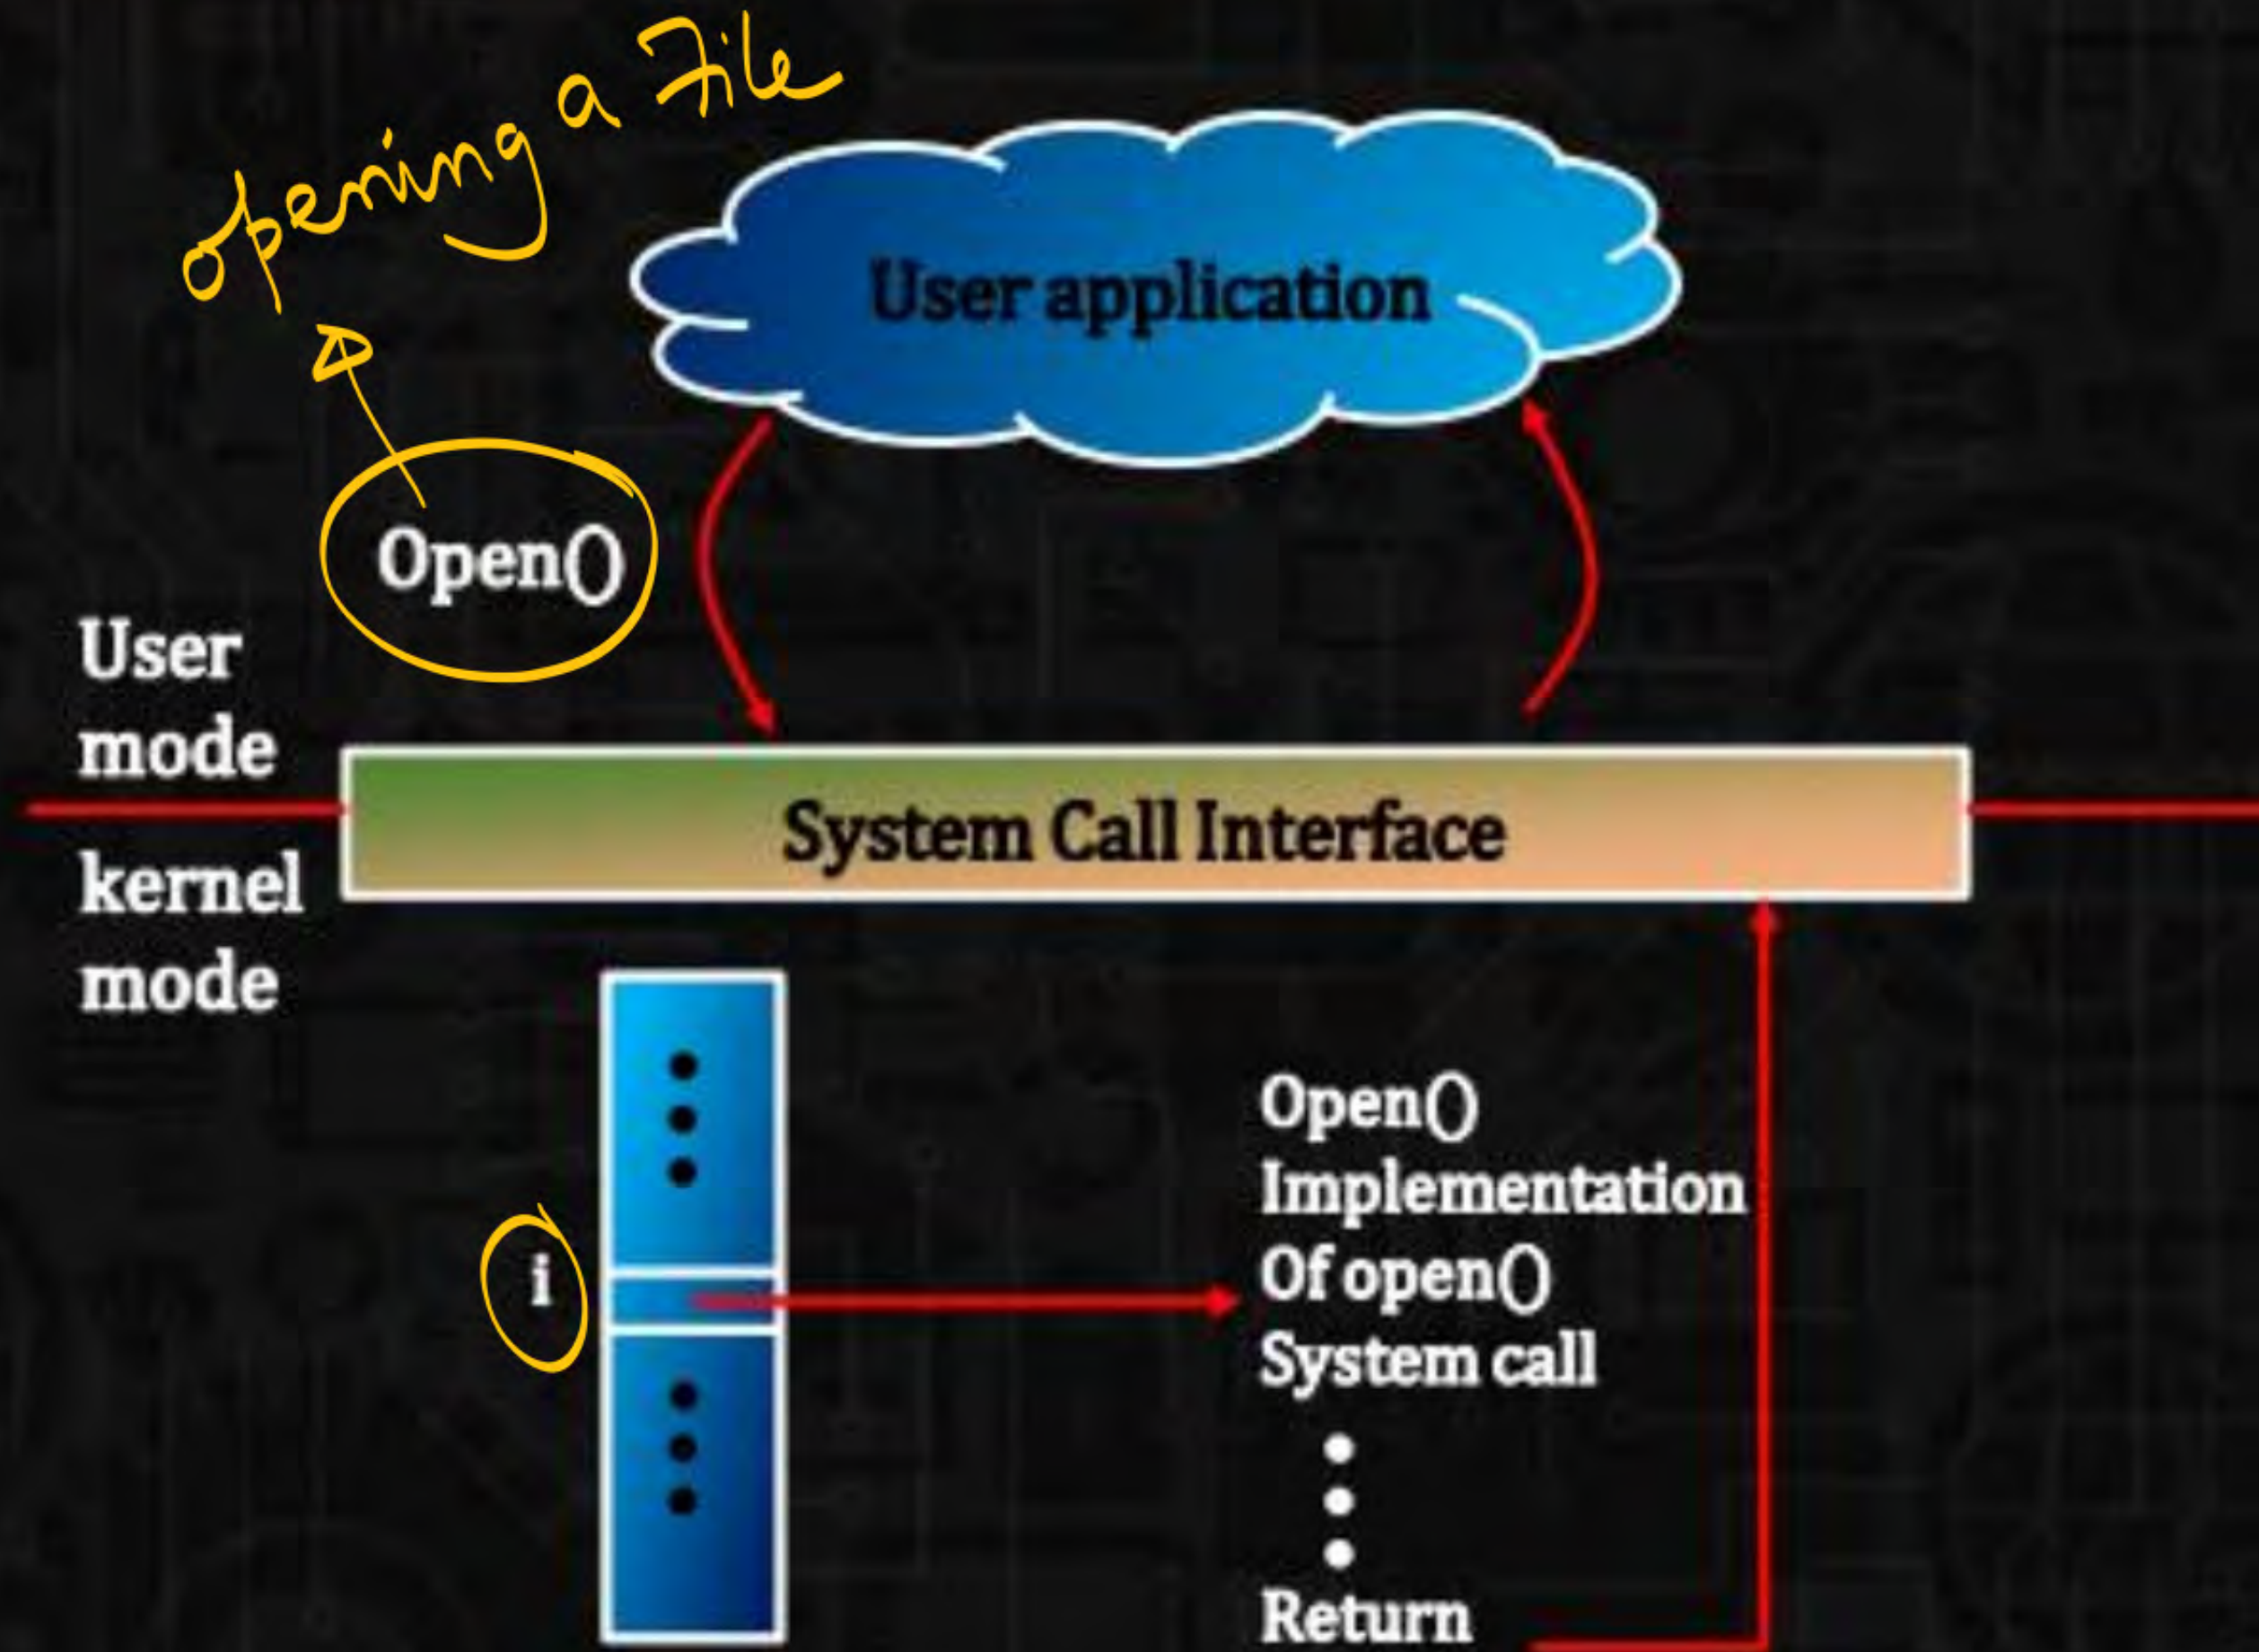
- ❑ Programming interface to the services provided by the OS
- ❑ Typically written in a high-level language (C or C++)
- ❑ Mostly accessed by programs via a high-level Application Program Interface (API) rather than direct system call use
- ❑ Three most common APIs are Win32 API for Windows, POSIX API for POSIX-based systems (including virtually all versions of UNIX, Linux, and Mac OS X), and Java API for the Java virtual machine (JVM)
- ❑ Why use APIs rather than system calls?  
(Note that the system-call names used throughout this text are generic)



- ❑ Typically, a number associated with each system call
  - ❖ System-call interface maintains a table indexed according to these numbers
- ❑ The system call interface invokes intended system call in OS kernel and returns status of the system call and any return values
- ❑ The caller need know nothing about how the system call is implemented
  - ❖ Just needs to obey API and understand what OS will do as a result call
  - ❖ Most details of OS interface hidden from programmer by API
- ❑ Managed by run-time support library (set of functions built into libraries included with compiler)



# API - System Call - OS Relationship

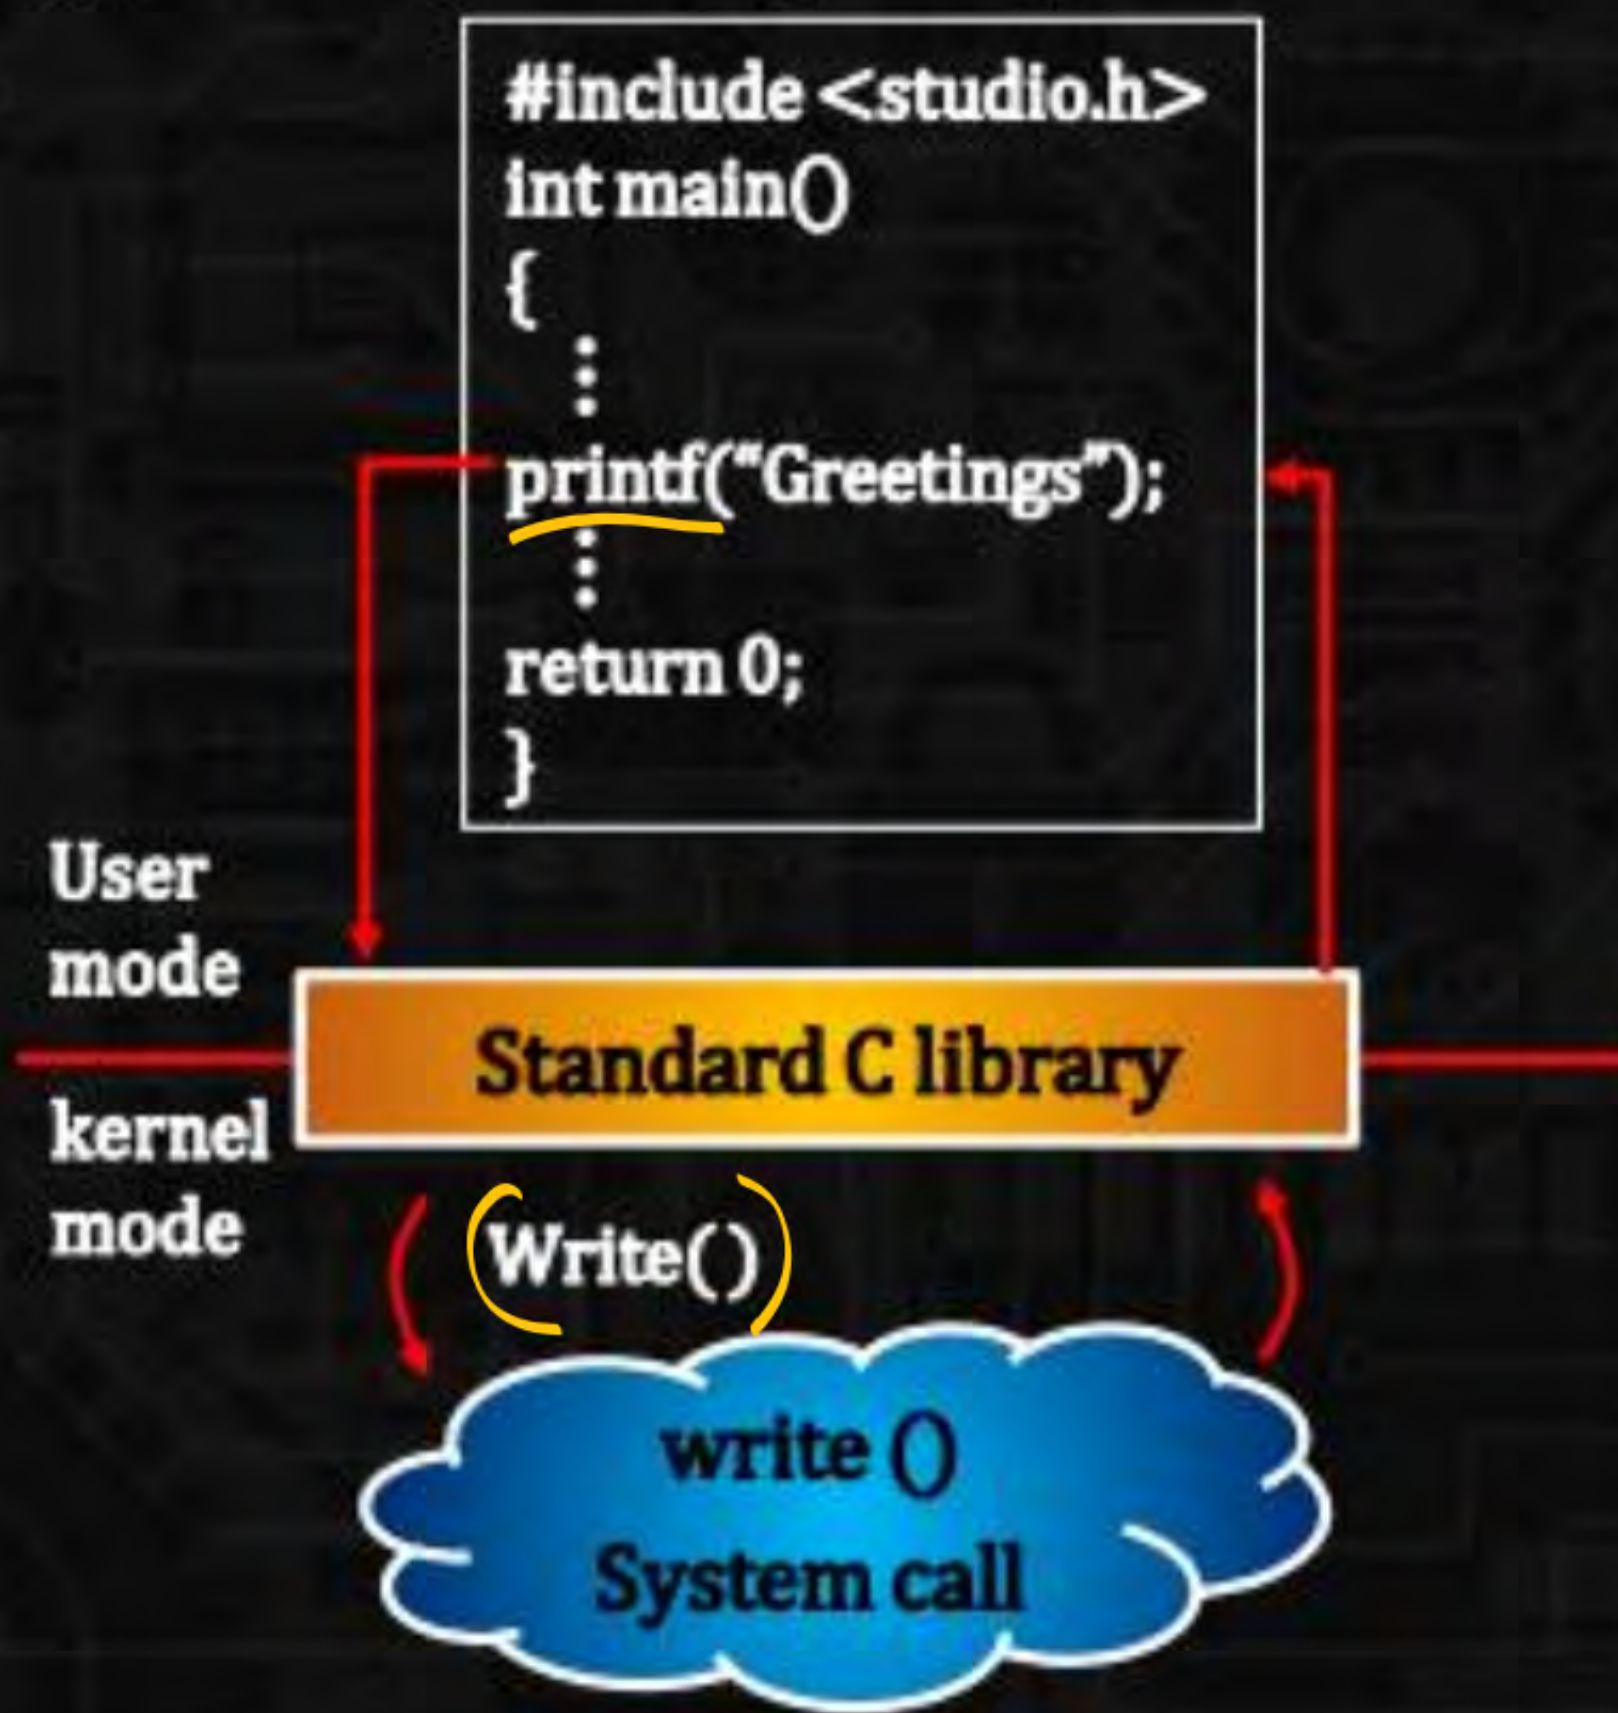




# Standard C Library Example

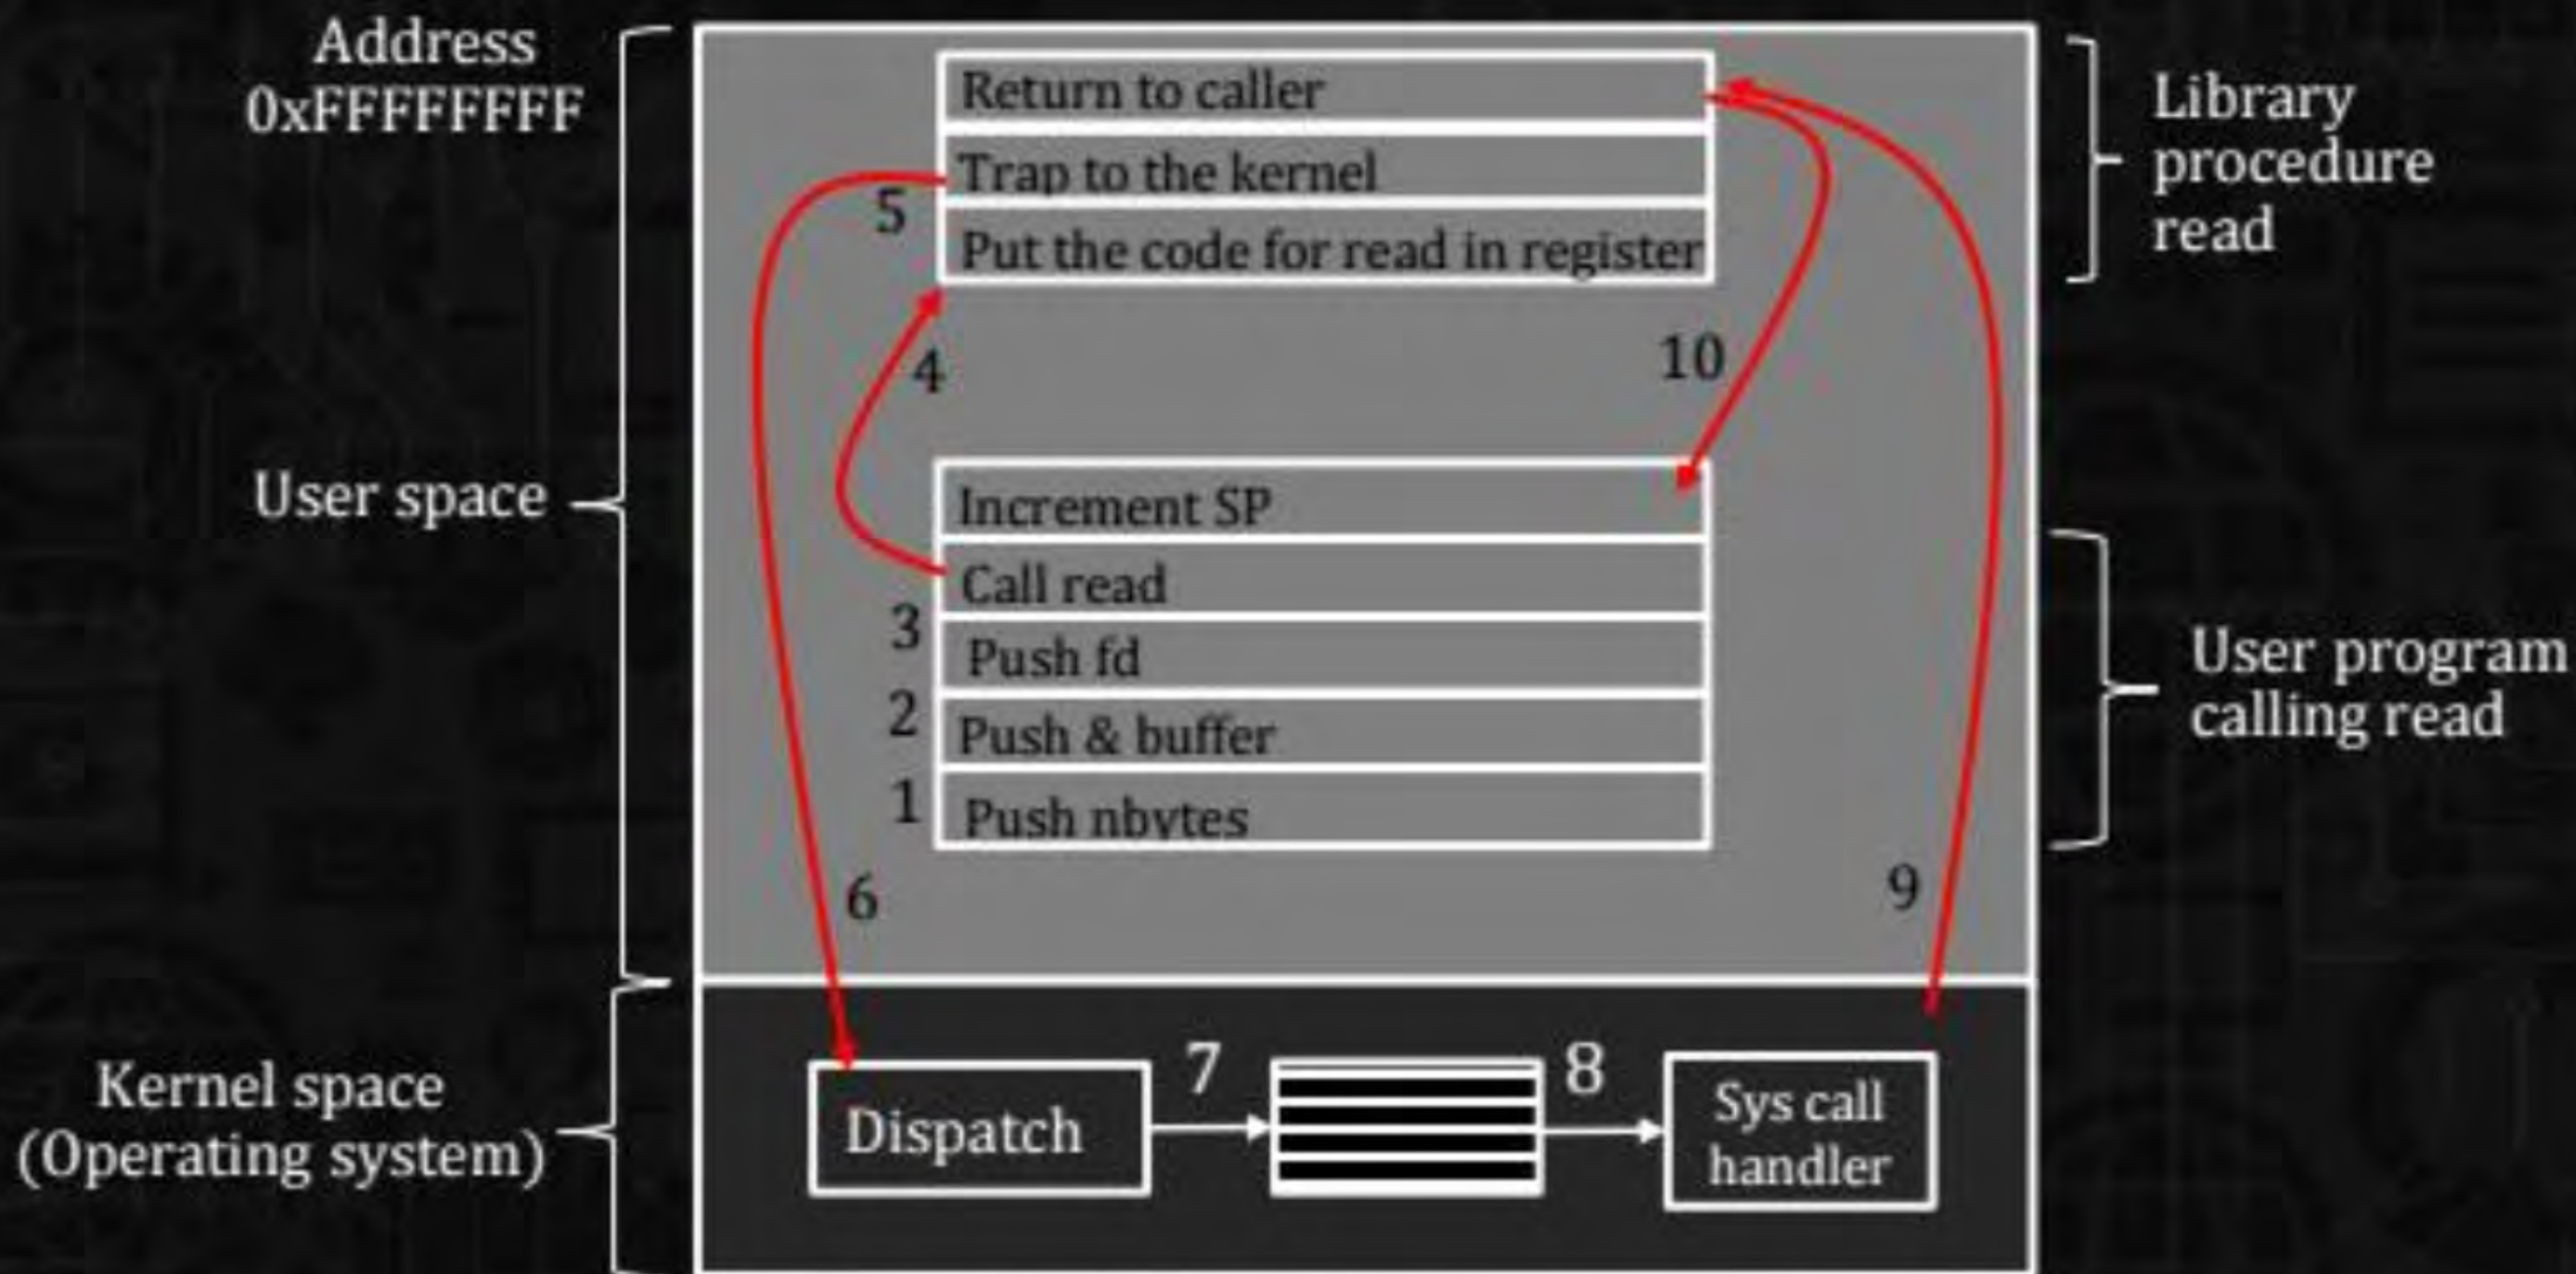


- ❑ C program invoking printf() library call, which calls write() system call





# Steps in Making a System Call



There are 11 steps in making the system call read (fd, buffer, nbytes)



# System Call Parameter Passing



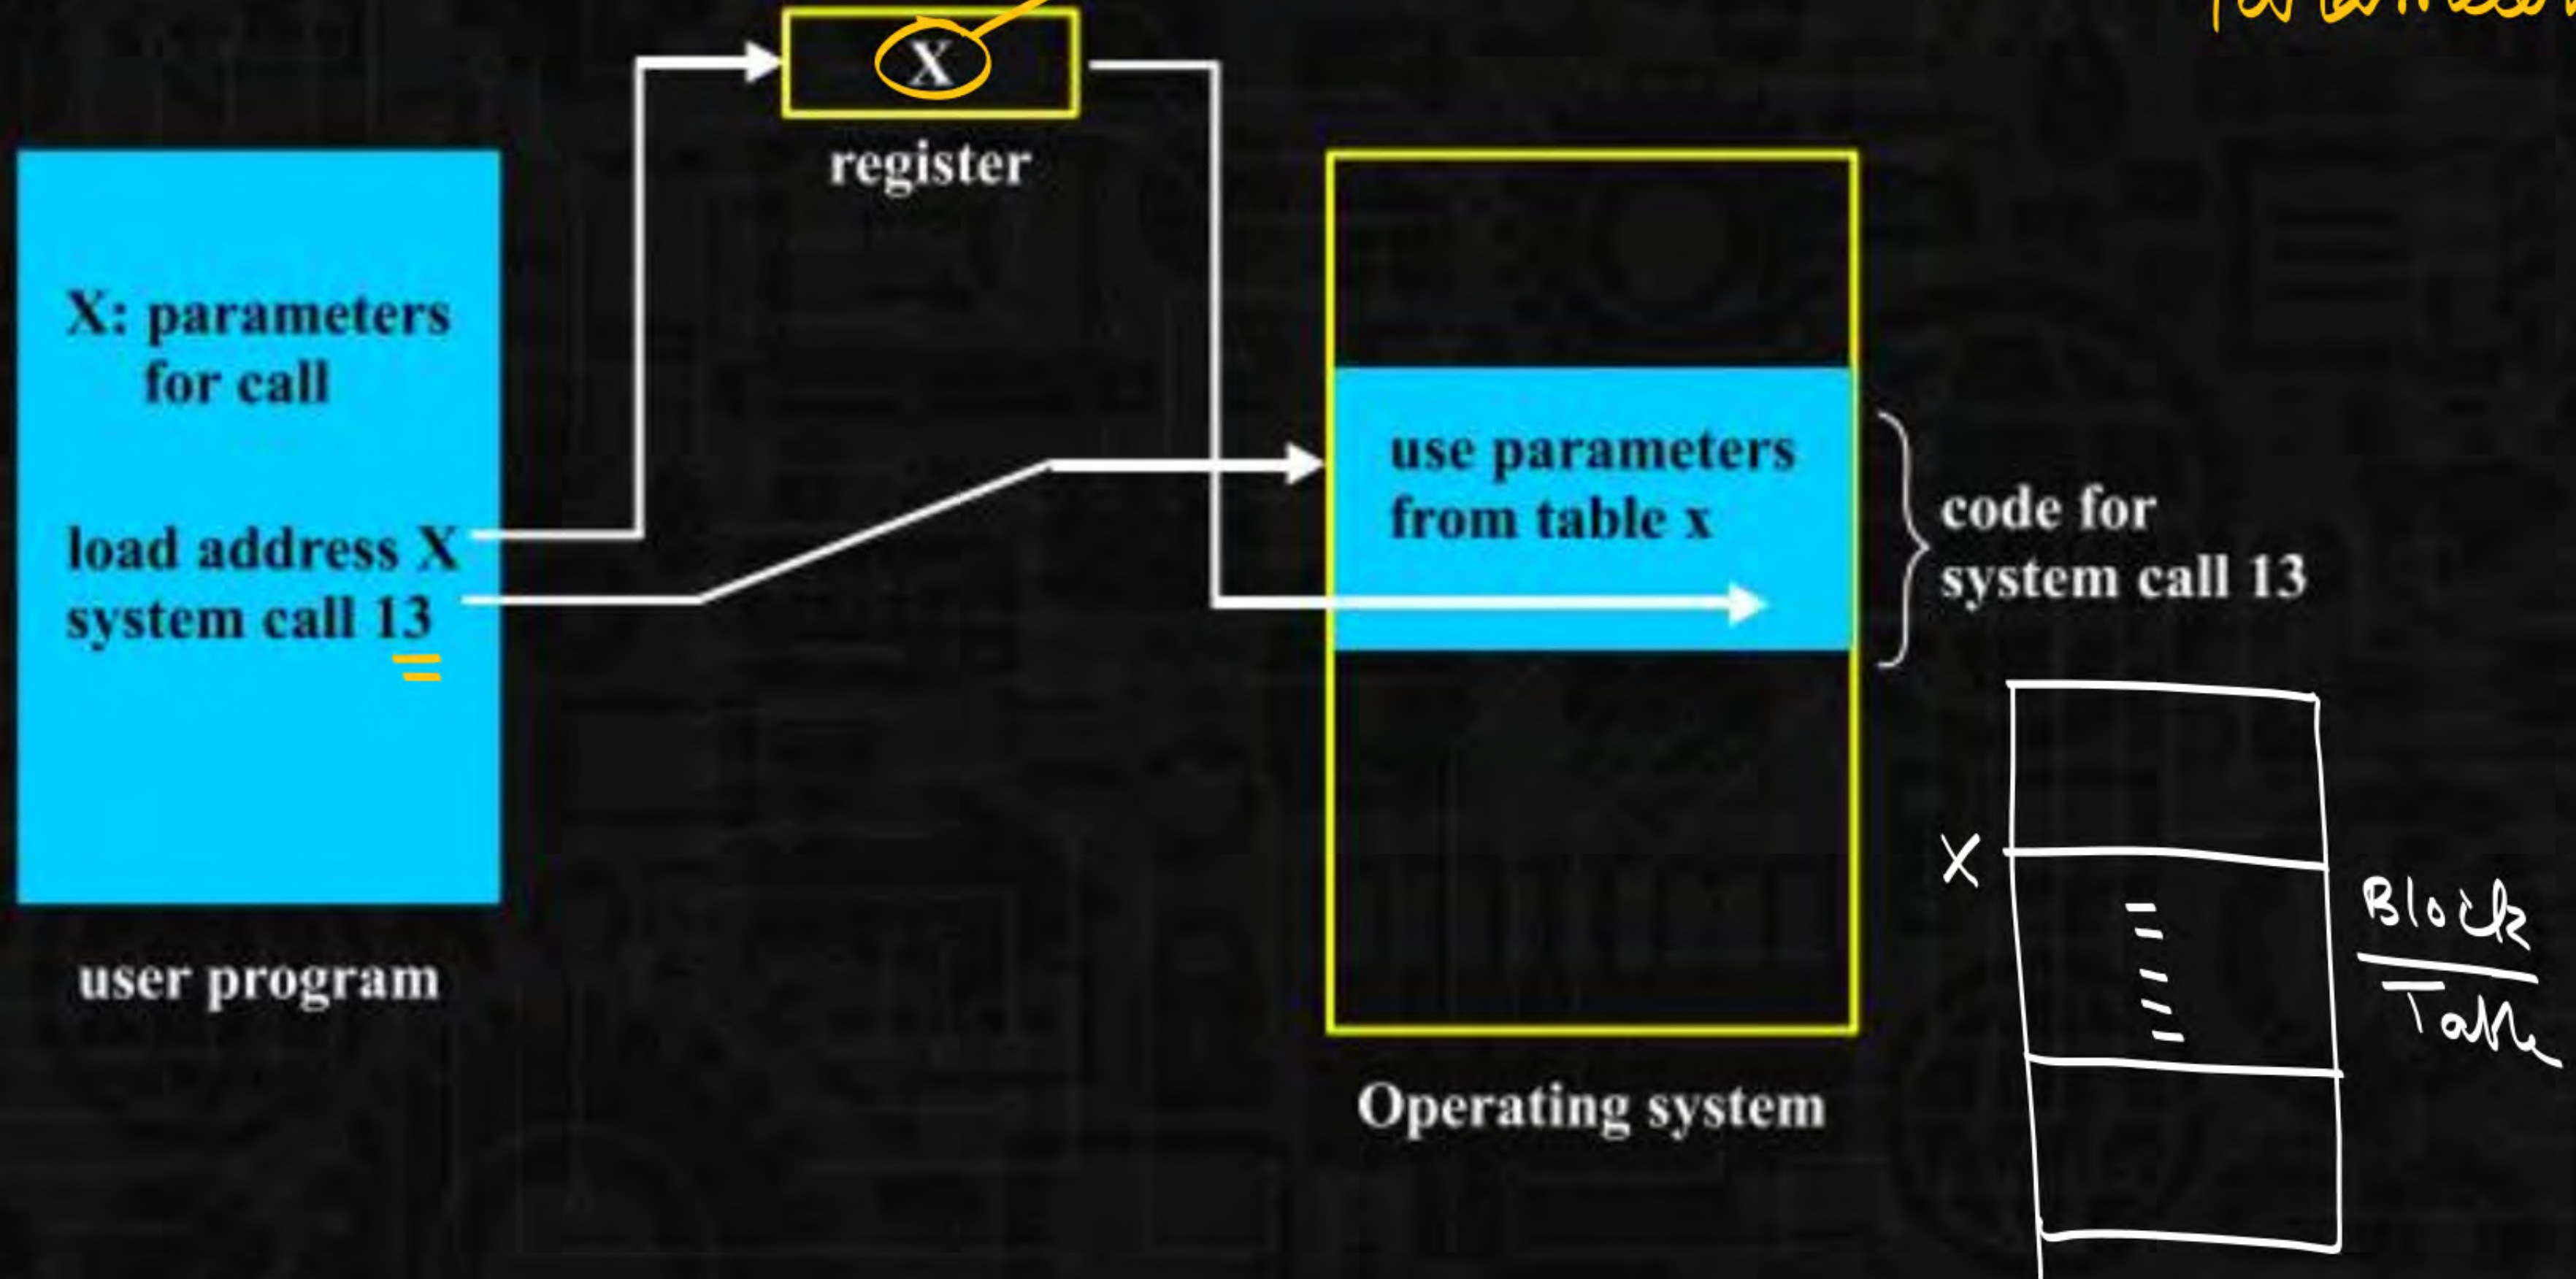
- ❑ Often, more information is required than simply identity of desired system call  
Exact type and amount of information vary according to OS and call
- ❑ Three general methods used to pass parameters to the OS
  - ❖ Simplest: pass the parameters in registers
    - In some cases, may be more parameters than registers
  - ❖ Parameters stored in a block, or table, in memory, and address of block passed as a parameter in a register
    - This approach taken by Linux and Solaris
  - ❖ Parameters placed, or pushed, onto the stack by the program and popped off the stack by the operating system  
Kernel
  - ❖ Block and stack methods do not limit the number or length of parameters being passed



# Parameter Passing via Table



Address of Block of Mem. containing Parameters;





## Types of System Calls

- ❑ Process control ✓
- ❑ File management ✓
- ❑ Device management ✓
- ❑ Information maintenance ✓
- ❑ Communications ✓



# Types of System Calls



## ❑ File management

- ❖ create file, delete file
- ❖ open, close file
- ❖ read, write, reposition
- ❖ get and set file attributes

## ❑ Device management

- ❖ request device, release device
- ❖ read, write, reposition
- ❖ get device attributes, set device attributes
- ❖ logically attach or detach devices



### ❑ Process control

- ❖ end, abort
- ❖ load, execute
- ❖ create process, terminate process
- ❖ get process attributes, set process attributes
- ❖ wait for time
- ❖ wait event, signal event
- ❖ allocate and free memory
- ❖ Dump memory if error
- ❖ Debugger for determining bugs, single step execution
- ❖ Locks for managing access to shared data between processes



# Types of System Calls



## ❑ Information maintenance

- ❖ get time or date, set time or date
- ❖ get system data, set system data
- ❖ get and set process, file, or device attributes

## ❑ Communications *<IPC Mechanisms>*

- ❖ create, delete communication connection
- ❖ send, receive messages if message passing model to host name or process name
  - From client to server
- ❖ Shared-memory model create and gain access to memory regions
- ❖ transfer status information
- ❖ attach and detach remote devices



## Types of System Calls

- ❑ Protection
  - ❖ Control access to resources
  - ❖ Get and set permissions
  - ❖ Allow and deny user access



## Some System Calls For Process Management







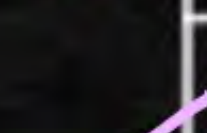


### Process Management

	Call	Description
①	<code>pid = fork()</code>	Create a child process identical to the parent
②	<code>pid = <u>waitpid</u>(pid, &amp;statloc, options)</code>	Wait for a child to terminate
③	<code>s = <u>execve</u>(name, argv, environp)</code>	Replace a process' core image
④	<code>exit(status)</code>	Terminate process execution and return status




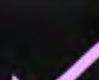
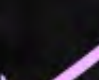


## Some System Calls For File Management



File Management	
Call	Description
 <code>fd = open(file, how, ...)</code>	Open a file for reading, writing or both
 <code>s = close(fd)</code>	Close an open file
 <code>n = read(fd, buffer, nbytes)</code>	Read data from a file into a buffer
 <code>n = write(fd, buffer, nbytes)</code>	Write data from a buffer into a file
 <code>position = lseek(fd, offset, whence)</code>	Move the file pointer
 <code>s = stat(name, &amp;buf)</code> 	Get a file's status information



## Some System Calls For Directory Management

Call	Description
 <code>s = mkdir(name, mode)</code>	Create a new directory
 <code>s = rmdir(name)</code>	Remove an empty directory
 <code>s = link(name1, name2)</code>	Create a new entry, name2, pointing to name1
 <code>s = unlink(name)</code>	Remove a directory entry
 <code>s = mount(special, name, flag)</code>	Mount a file system
<code>s = umount(special)</code>	Unmount a file system



## Some System Calls For Miscellaneous Tasks



Miscellaneous	
Call	Description
<code>s = chdir(dirname)</code>	Change the working directory
<code>s = chmod(name, mode)</code>	Change a file's protection bits
<code>s = kill(pid, signal)</code>	Send a signal to a process
<code>seconds = time(&amp;seconds)</code>	Get the elapsed time since Jan 1,1970



# Examples of Unix and Windows System Calls



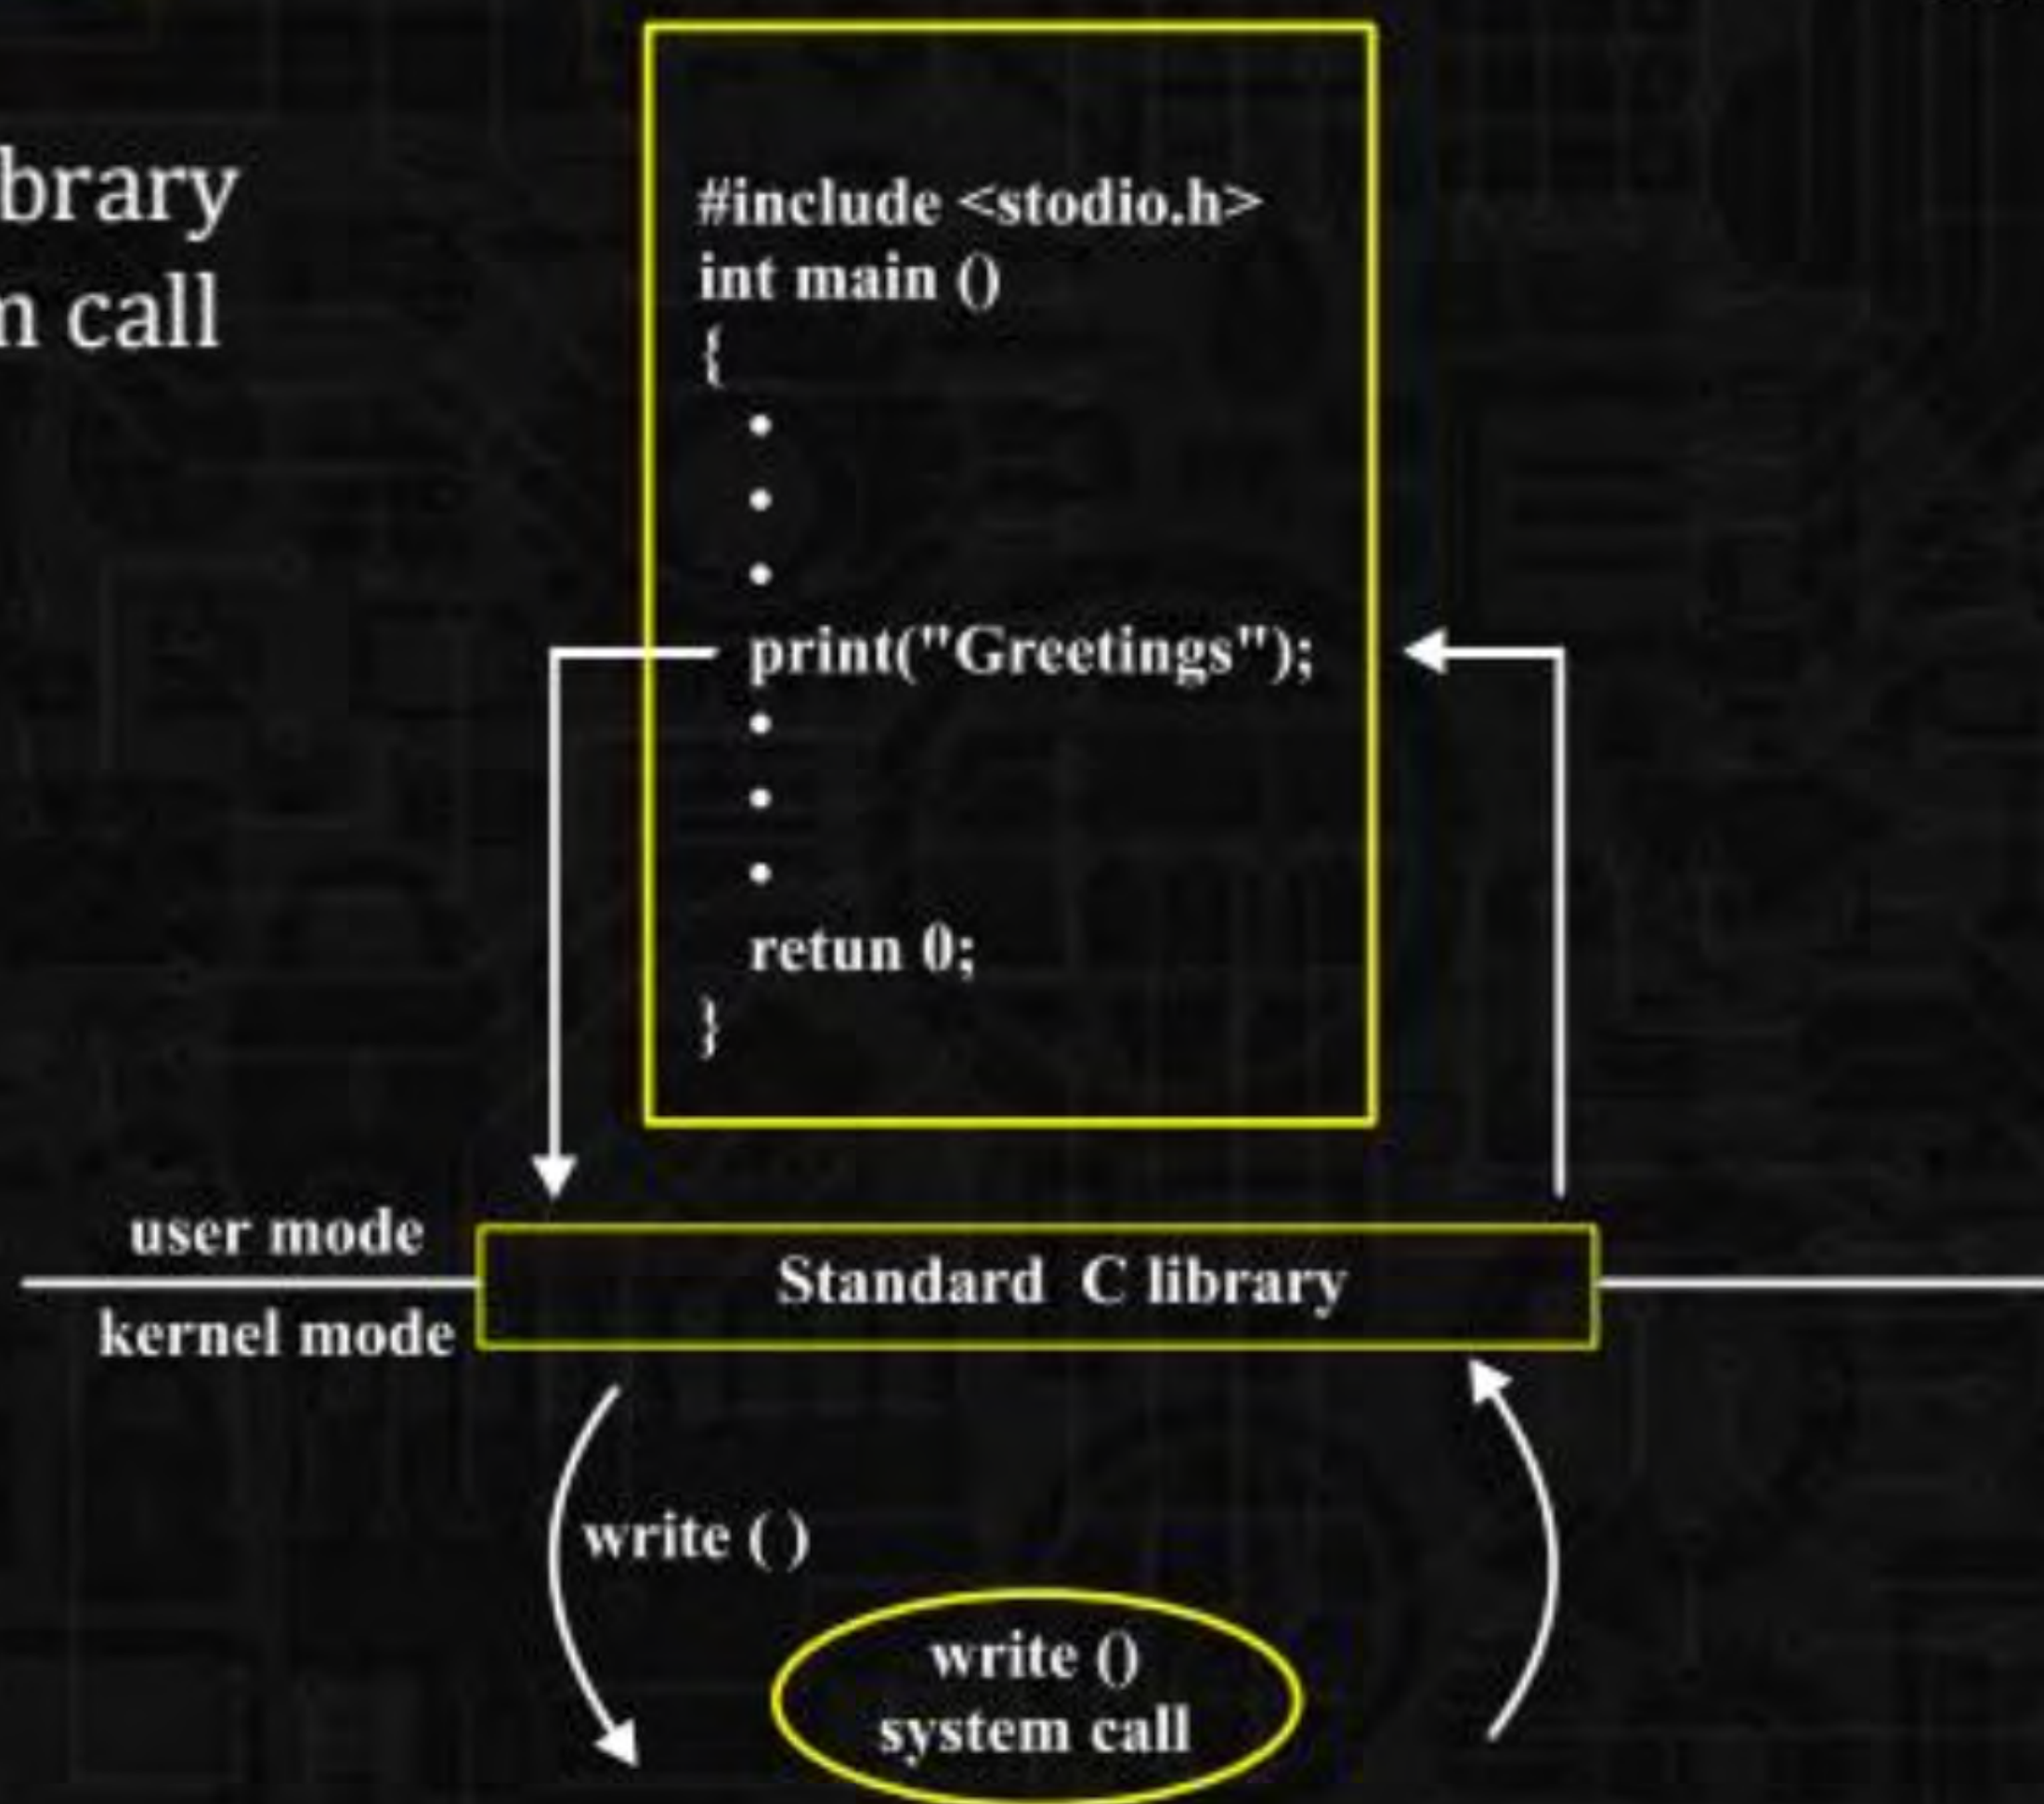
	Windows	Unix
<b>Process control</b>	CreateProcess() ExitProcess() WaitForSingleObject()	fork() exit() wait()
<b>File Manipulation</b>	CreateFile() ReadFile() WriteFile() CloseHandle()	open() read() write() close()
<b>Device manipulation</b>	SetConsoleMode() ReadConsole() WriteConsole()	ioctl() read() write()
<b>Information maintenance</b>	GetCurrentProcessID() SetTimer() Sleep()	getpid() alarm() sleep()
<b>Communication</b>	CreatePipe() CreateFileMapping() MapViewOfFile()	pipe() shmget() mmap()
<b>Protection</b>	SetFileSecurity() InitializeSecurityDescriptor() SetSecurityDescriptorGroup()	chmod() umask() chown()



## Standard C Library Example



- ❑ C program invoking printf() library call, which calls write() system call





## Common and well known system calls are:

- ❑ access: checks if calling process has file access
- ❑ alarm: sets a process's alarm clock
- ❑ chdir: changes the working directory
- ❑ chmod: changes the mode of a file
- ❑ chown: changes the ownership of a file
- ❑ chroot: changes the root directory
- ❑ close: closes a file descriptor
- ❑ dup, dup2: duplicates an open file descriptor



## Common and well known system calls are:

- ❑ execl, execv, execl, execve, execlp, execvp: executes a file
- ❑ exit: exits a process
- ❑ fcntl: controls open files
- ❑ fork: creates a new process
- ❑ getpid, getpgrp, getppid: gets group and process IDs
- ❑ getuid, geteuid, getgid, getegid: gets user and group IDs
- ❑ ioctl: controls character devices
- ❑ kill: sends a signal to one or more processes
- ❑ link: links a new file name to an existing file
- ❑ lseek: moves read/write file pointer



## Common and well known system calls are:

- ❑ `mknod`: makes a directory, special or ordinary file
- ❑ `mount`: mounts a filesystem
- ❑ `msgctl`, `msgget`, `msgsnd`, `msgrcv`: message passing support
- ❑ `nice`: changes priority of a process
- ❑ `open`: opens a file for reading or writing
- ❑ `pipe`: creates an interprocess pipe
- ❑ `plock`: locks a process in memory
- ❑ `ptrace`: allows a process to trace the execution of another



## Common and well known system calls are:



- ❑ read: reads from a file
- ❑ semctl, semget, semop: semaphore support
- ❑ setpgrp: sets process group ID
- ❑ setuid, setgid: sets user and group IDs
- ❑ shmctl, shmget, shmop: shared memory support
- ❑ signal: control of signal processing
- ❑ sleep: suspends execution for an interval
- ❑ stat, fstat: gets file status
- ❑ stime: sets the time
- ❑ sync: updates the super block



## Common and well known system calls are:



- ❑ time: number of seconds since 1/1/1970
- ❑ times: gets process and child process times
- ❑ ulimit: gets and sets user limits
- ❑ umask: gets and sets file creation mask
- ❑ umount: unmounts a file system
- ❑ uname: gets system information
- ❑ unlink: removes directory entry
- ❑ ustat: gets file system statistics
- ❑ utime: sets file access and modification times
- ❑ wait: waits for a child process to stop or terminate
- ❑ write: writes to a file



- ❑ Provide a convenient environment for program development and execution
  - ❖ Some of them are simply user interfaces to system calls; others are considerably more complex
- ❑ File management - Create, delete, copy, rename, print, dump, list, and generally manipulate files and directories
- ❑ Status information
  - ❖ Some ask the system for info - date, time, amount of available memory, disk space, number of users
  - ❖ Others provide detailed performance, logging, and debugging information
  - ❖ Typically, these programs format and print the output to the terminal or other output devices
  - ❖ Some systems implement a registry - used to store and retrieve configuration information



- ❑ **File modification**
  - ❖ Text editors to create and modify files
  - ❖ Special commands to search contents of files or perform transformations of the text
- ❑ **Programming-language support** - Compilers, assemblers, debuggers and interpreters sometimes provided
- ❑ **Program loading and execution** - Absolute loaders, relocatable loaders, linkage editors, and overlay-loaders, debugging systems for higher-level and machine language
- ❑ **Communications** - Provide the mechanism for creating virtual connections among processes, users, and computer systems
  - ❖ Allow users to send messages to one another's screens, browse web pages, send electronic-mail messages, log in remotely, transfer files from one machine to another



## Fork System Call

### Case Study:

- create a child Process;
- The code of child will be an exact Replica/Copy of Parent Process;
- Execution in child will start from next Stmt. after fork, till end of Program;




```
main()  
{  
  printf("Hello");  
}
```

Process



Hello

```
main()  
{  
  1. fork();  
  2. printf("Hello");  
}
```






```

main()
{
1. printf("Hi");
2. fork(); ✓
3. printf("Hello");
}

```

Total: 2  
child: 1




```

main()
{
1. fork();
2. fork();
3. printf("Hello");
}

```

Total Processes: 4  
Child Processes: 3  
New





```
main()
```

```
{  
1. fork();
```

```
2. fork();
```

```
3. fork();
```


```
4. printf("Hello");
```

```
}
```

3 forks:

Total : 8 ✓

child : 7 ✓





No. of forks	Total Processes	Child Processes
1	2	1
2	4	3
3	8	7
4	16	15
'n'	$2^n$	$(2^n - 1)$



```
main()  
{
```

```
1. printf("one");
```

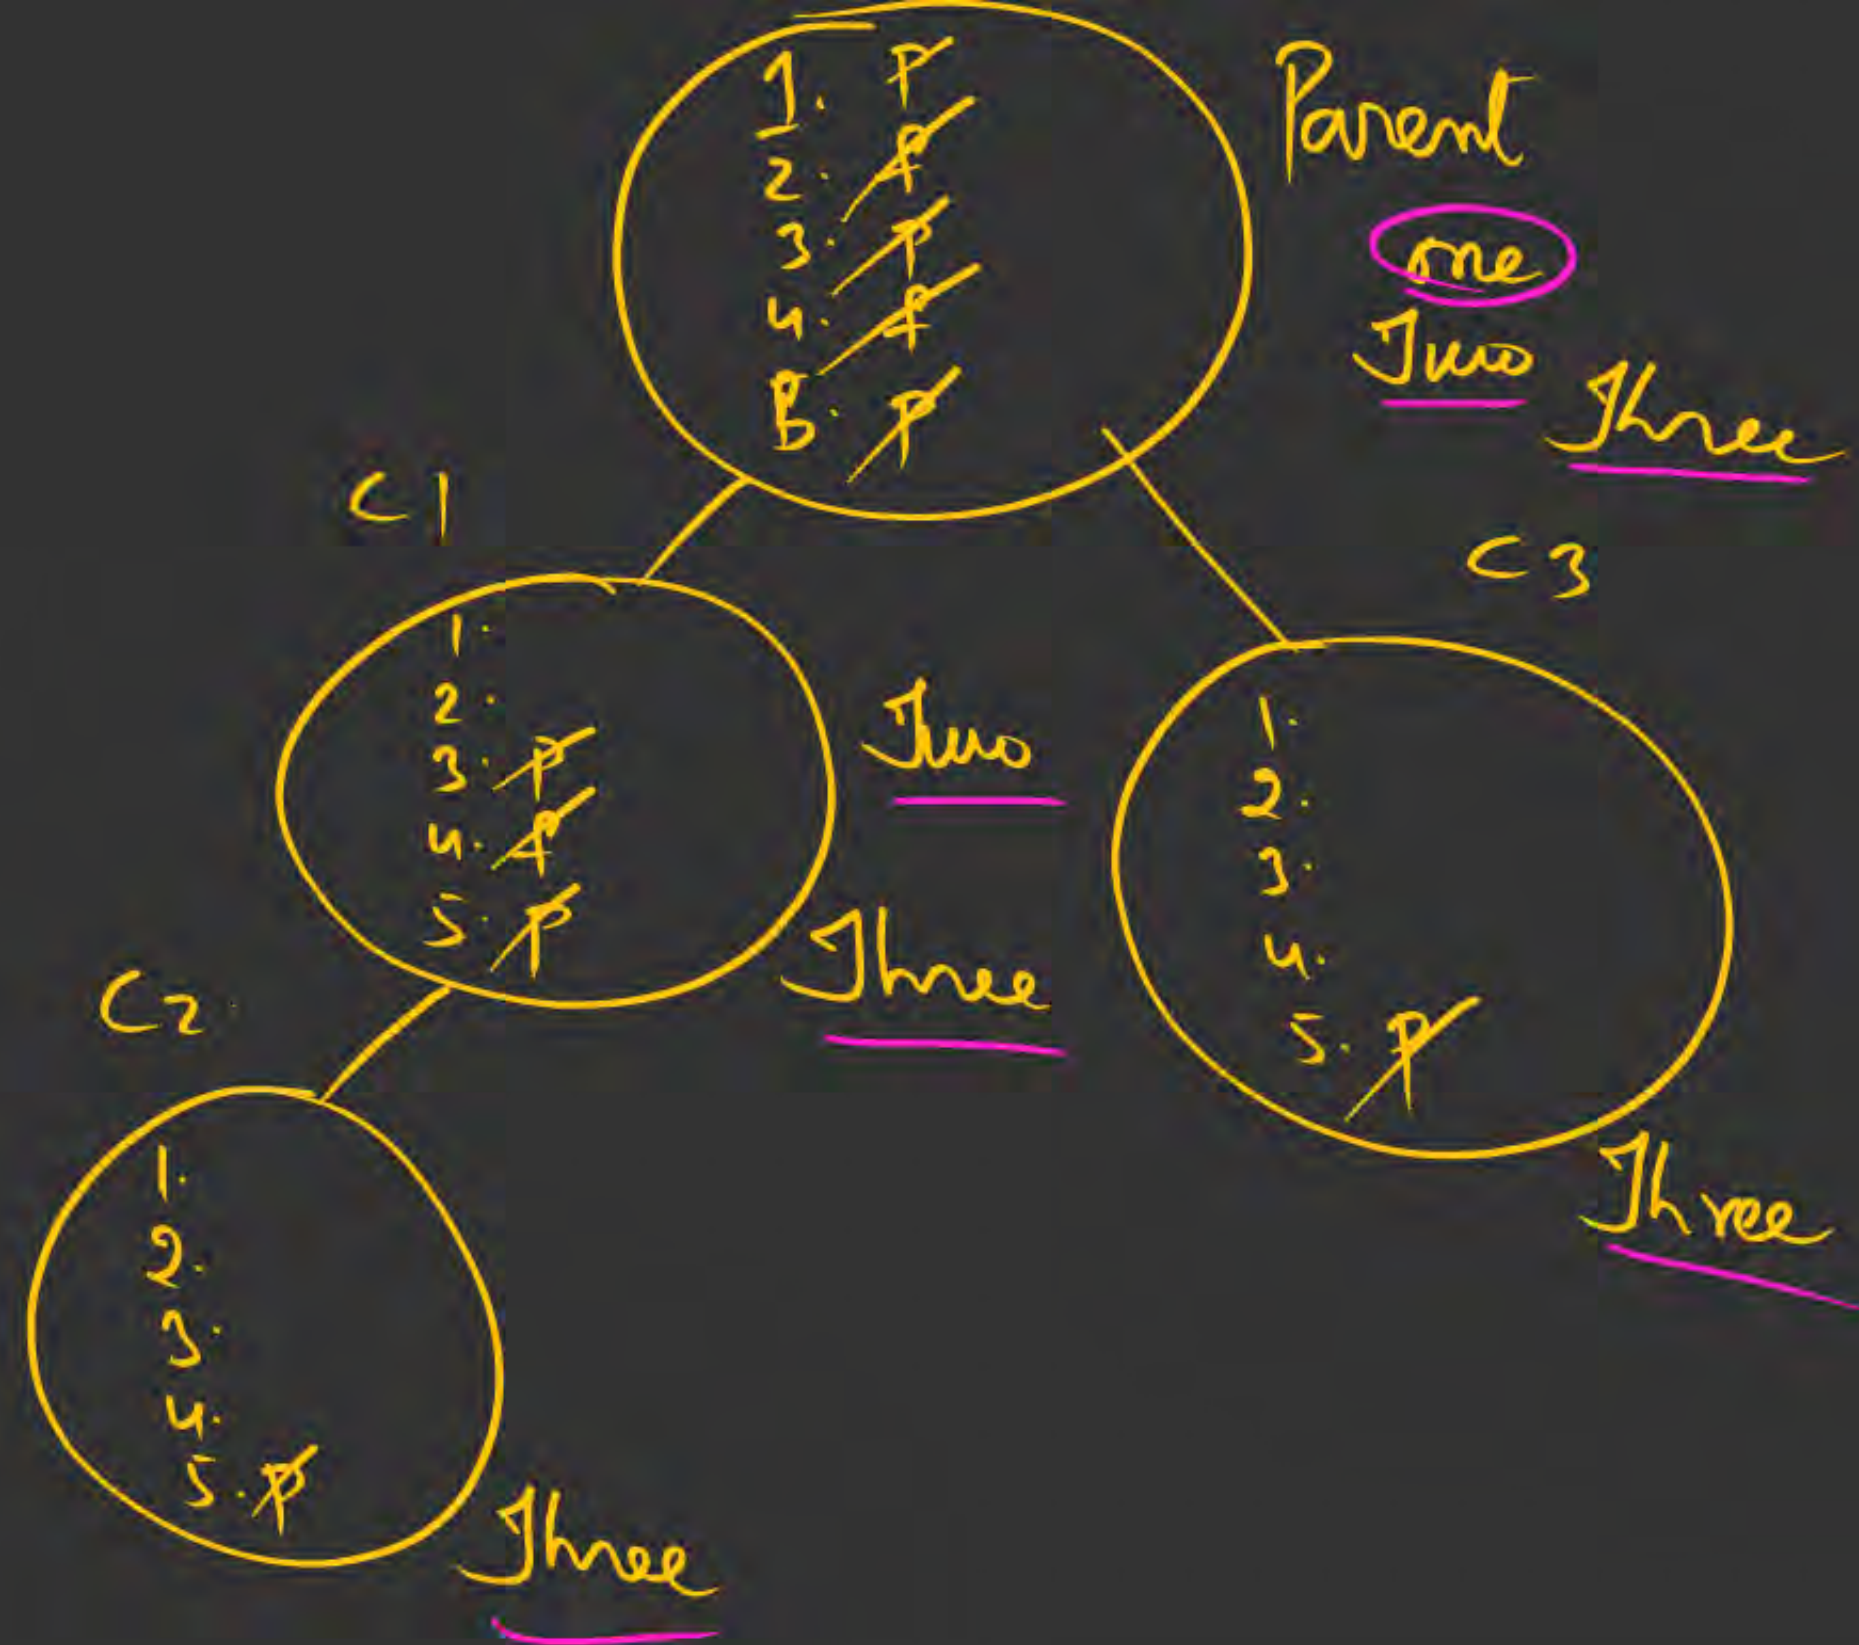
```
2. fork();
```

```
3. printf("Two");
```

```
4. fork();
```

```
5. printf("Three");
```

```
}
```





```

main()
{
    int i, n;

    for(i=1; i<=n; ++i)
        fork();
}

```

How many child Processes  
are created =

```

fork();
fork();
fork();
⋮
fork();

```


} 'n' Times

$= 2^n$

Child:  $(2^n - 1)$




Return value of fork()



{ Assume: fork() always Succeed }

```
main()
{
    int ret;
    ret = fork() <Parent>;
    if (ret == 0)
    {
        // child //
    }
    else
    {
        // Parent
    }
    == <child + Parent>
}
```





main()

{

int a=5, b=3, c;

c = ++a \* b++;

Pf(a, b, c);


if (fork() == 0)

{  
a += 5;  
b += 3;  
printf(a, b);  
c--;

}  
else

{  
b += a + 2;  
Pf(a, b);  
c--;


}  
Pf(c);



→ 6, 4, 18

→ 8, 12

→ 17



→ 11, 7

→ 17



# Q.1

```
main ()
{
```

pyq ⊗

```
    int i, n;
    for (i = 1; i <= n; ++i)
    if (fork () == 0);
        print ("*");
}
```

No. of times '\*' get printed =  $\frac{n}{2}$

→ every Process will execute

```
int i, n;
for (i = 1; i <= n; ++i)
{
    if ( $\frac{i}{2} == 0$ )
        (fork());
}
```

No. of children:  $(2^{\frac{n}{2}} - 1)$



Assume fork always succeeds

```
main()
{
  int a;
```

```
  if (fork() == 0)
  {
```

```
    a = a + 5;
    print(a); → u
```

```
  }
  else
  {
```

```
    a = a - 5;
    print(a); → x
```

```
  }
}
```

&a → Same(VA)

&a → Same(VA)

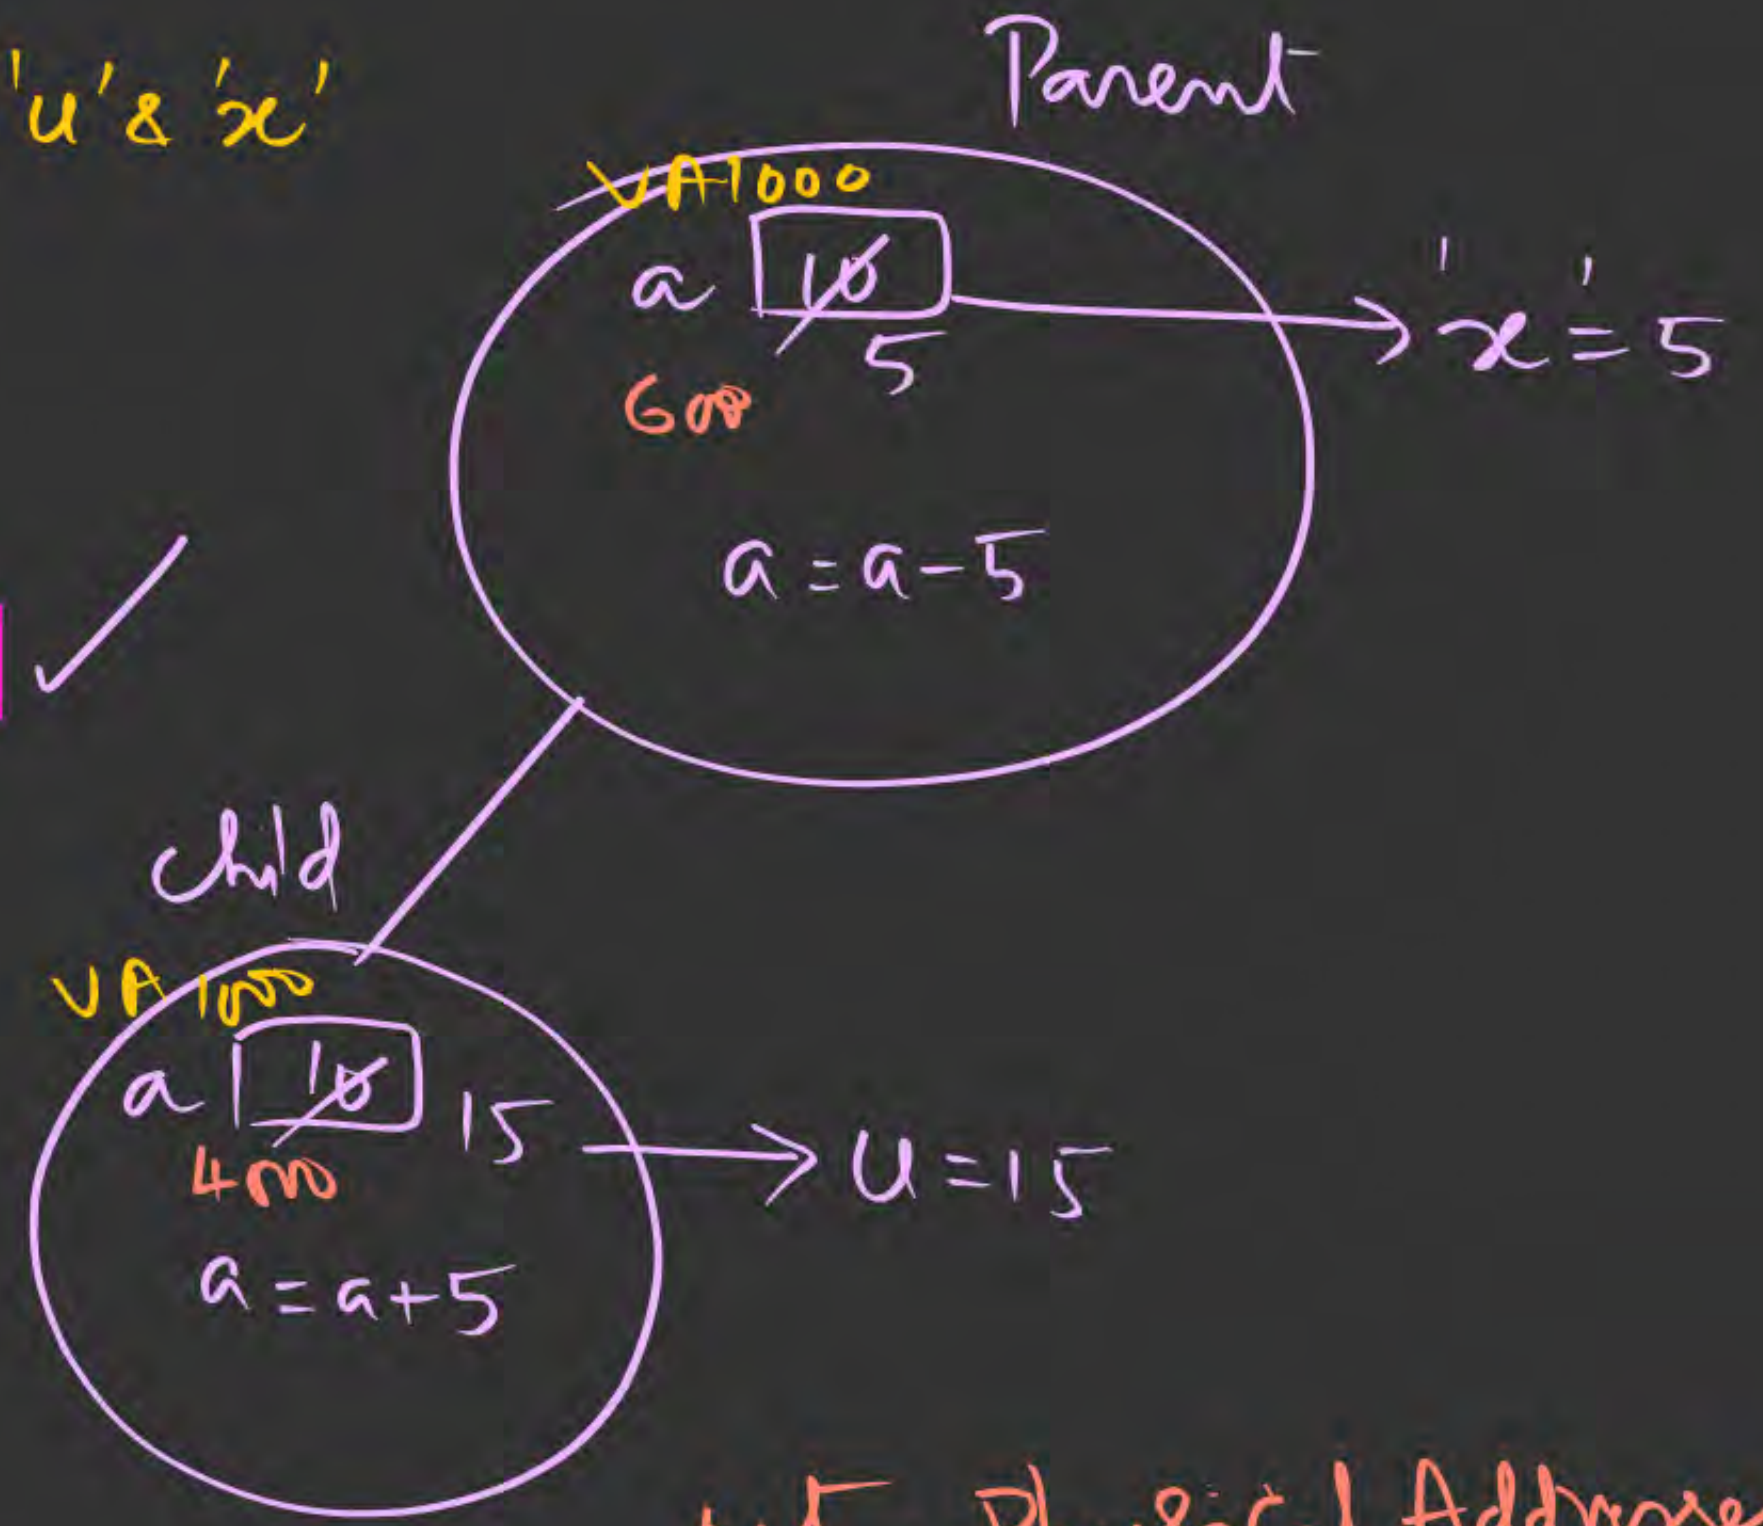
⇒ What is the Relation b/w 'u' & 'x'?

a)  $u = x$

b)  $u = x + 5$

c)  $u = x + 10$  ✓

d)  $x = u - 5$



Note: 1) Physical Addresses will be different

2) virtual Addresses of 'a' will be same



Q. 2

main ()

{

int i, n;

for (i = 1; i &lt;= n; ++i)

{

1. fork ();

2. print ("\*");

}

}

H/W



Q. 3

main () Q<sub>3</sub>  
{

H/w

```
int i, n;  
for (i = 1; i <= n; ++ i)  
{  
    1. print ("*");  
    2. fork ();  
}  
}
```

" \* "

formal

Q<sub>2</sub>

```
for i ← 1 to n  
{  
    fork();  
    print(*);  
}
```

" \* "

formule



Q. 3

main ()

{

int i, n;

for (i = 1; i &lt;= n; ++ i)

{

print ("\*");

fork ();

}

}



# Q. 4

```
main ()
{
```

```
    int a = 1, b = 2, c = 3;
    a += b += ++ c;
    print (a, b, c);
    if (fork () == 0)
    {
        int d;
        ++ a; ++ b; -- c;
        print (a, b, c);
        if (fork () == 0)
        {
            d = a + b + c;
            print (a, b, c, d);
        }
    }
```

```
else
```

```
{
    -- a; -- b;
    c = a + b; d = a + b + c;
    print (a, b, c, d);
}
}
else
{
    c += b += ++ a;
    print (a, b, c) ;
}
print (d);
```

```
}
```



**Q. 5**

The following C program is executed on a Unix/Linux system:

```
#include <unistd.h>
int main()
{   int i;
    for (i = 0; i < 10; i++)
        if (i % 2 == 0) fork();
}
```

The total number of child processes created is

31

P4Q



fork(); exec; wait

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

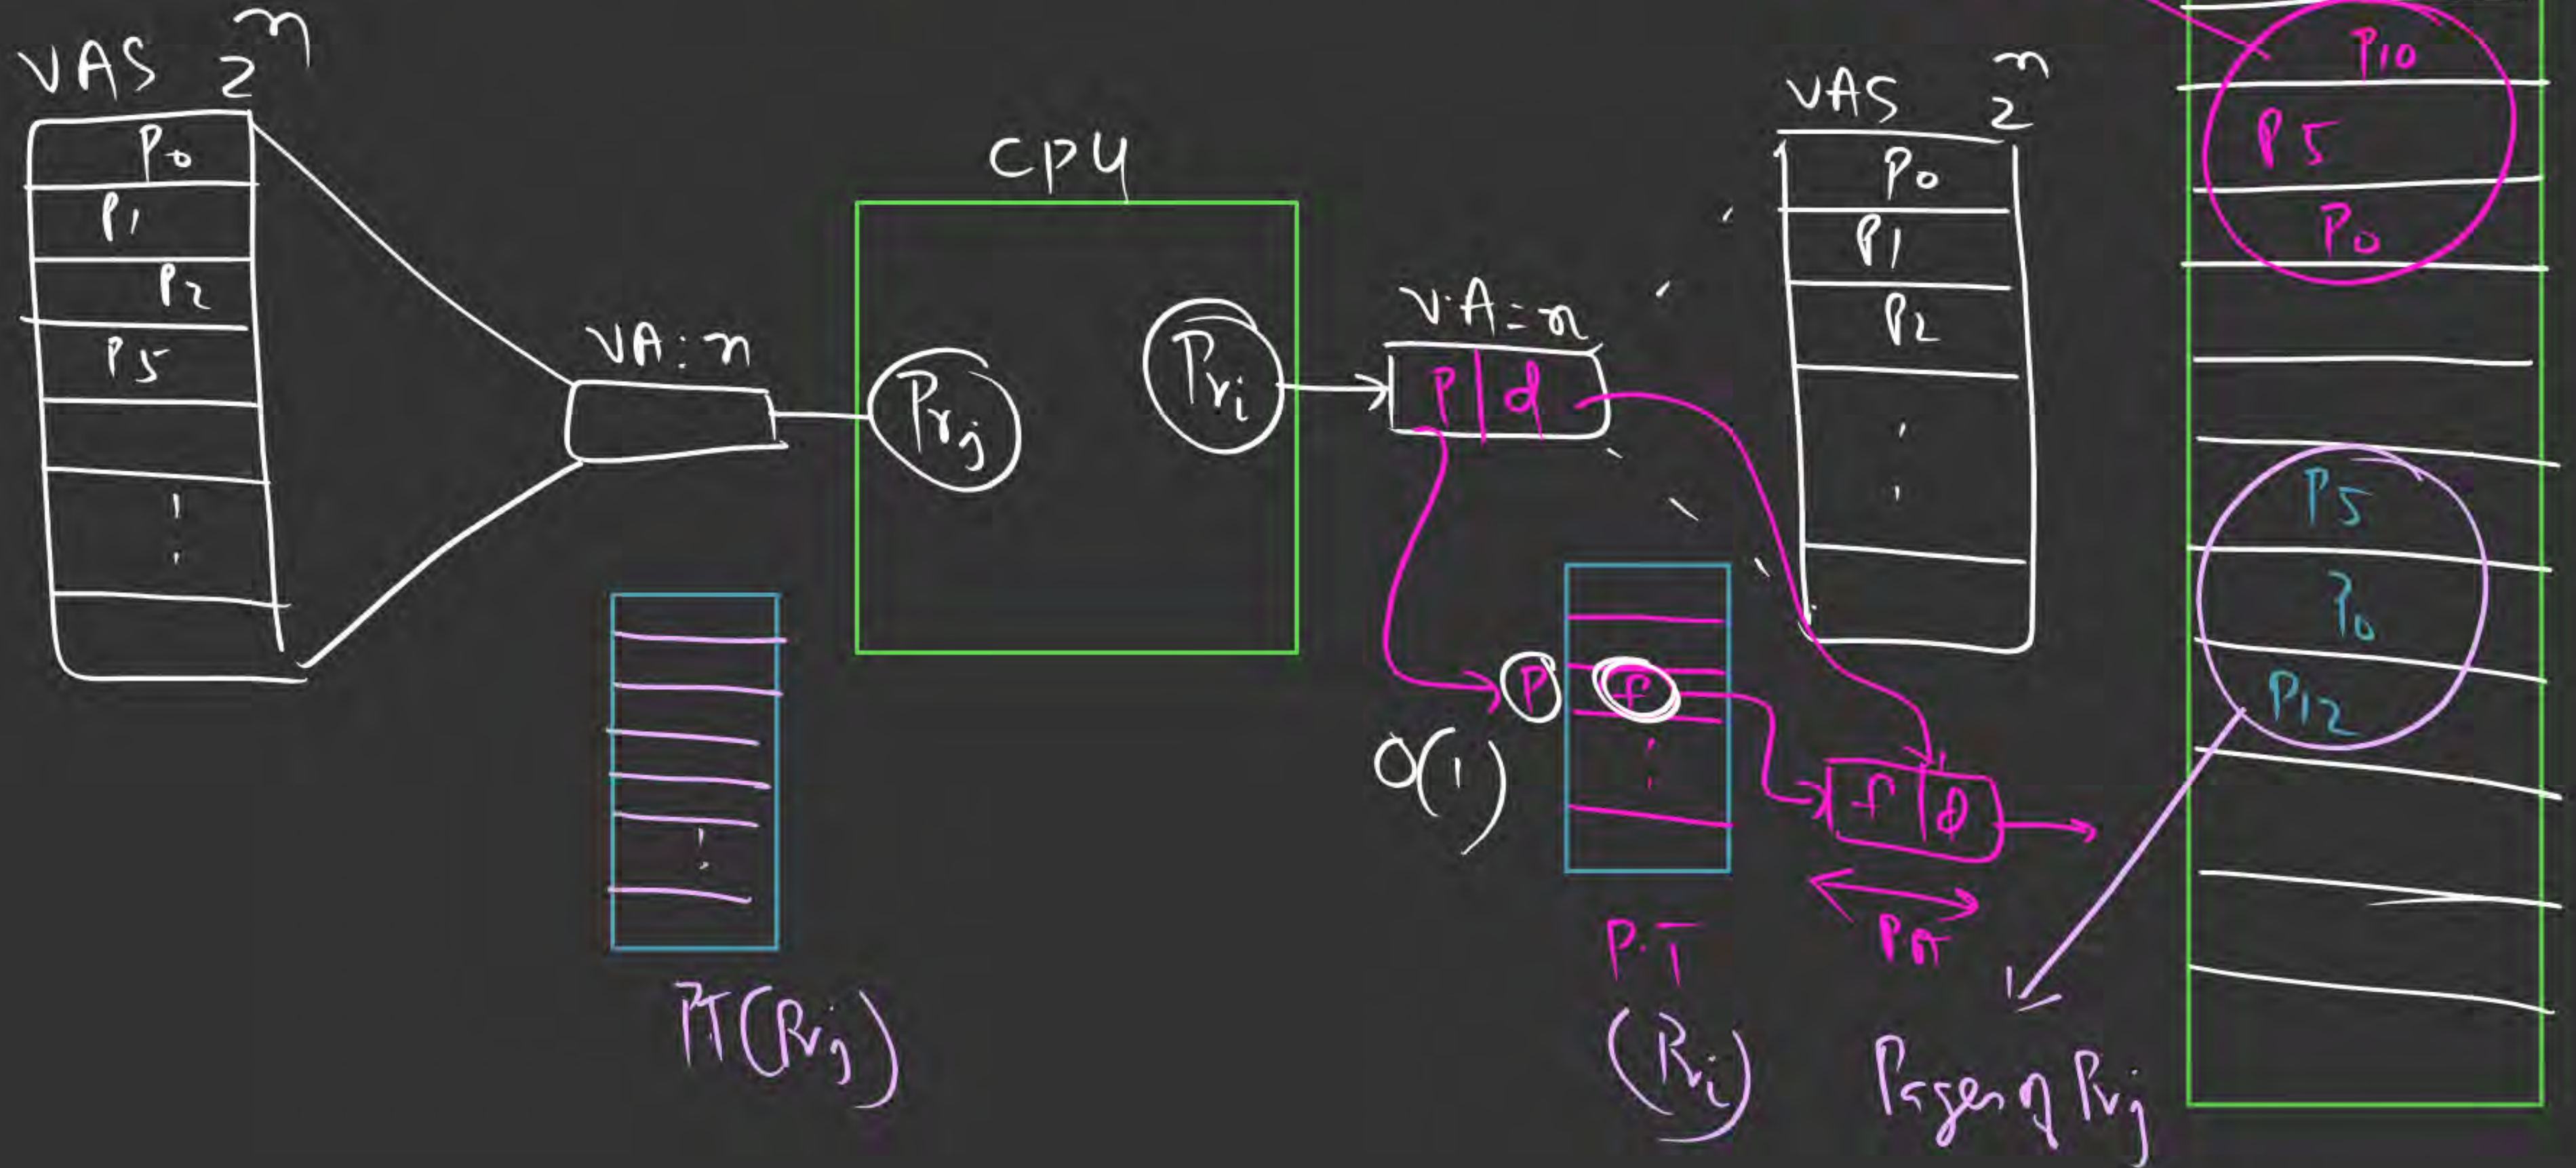
    /* fork a child process */
    pid = fork(); ✓
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed"); ✓
        return 1;
    } else if (pid == 0) { /* child process */
        execvp ("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    } return 0;
}
```



# Inverted - Paging


## Reverse Paging

< Reduce Space overhead of P.T in Paging >





In Inverted Paging, one global P.T is designed to be used by all Processes,





Consider a System with  $V.A = 34$  bits &  $P.A = 29$  bits;

$P.S = 8KB$ ;  $P.T.E = 32$  bits;

What is the Size of

(i) Inadditional P.T:  $\left( \frac{2^{34}}{2^{13}} \right) \times 2^2 B = 2^{21+2} = 2^{23} = \underline{\underline{8MB}}$  ✓

(ii) Inverted P.T:  $\left( \frac{2^{29}}{2^{13}} \right) \times 2^2 = 2^{16+2} = 2^{18} = \underline{\underline{256KB}}$



# Paging vs Segmentation

	I.F	E.F
Paging	✓ <Last Page>	✗
Segment	✗	✓ <P.A.S>

2) To reduce Seg-Table  
Size overhead

we apply Paging  
on S.T

1) Compaction

2) Paging on Segment

① <Segmented - Paging>



